

Paise Migration Audit Report

BahiKhata Pro — V17 Float-to-Paise Migration Initiative

Report Date	2026-07-12
Prepared By	AI Engineering Agent (Z.ai)
Audience	External Auditor / Technical Reviewer
Phases Complete	Phase 1, 2 (A-G), 3 — DONE
Phases Remaining	Phase 4, 5, 6, 7 — PENDING DECISION
Total Bugs Found	10 (5 fixed, 4 open, 1 data issue)
Regression Tests Added	54 across 7 sub-phases

1. Executive Summary

BahiKhata Pro is a financial ledger application for Indian shop owners. Like most applications handling currency, it stores money values as **Float** (IEEE 754 double-precision) in the database. This is a well-known source of precision bugs: $0.1 + 0.2 = 0.30000000000000004$ in JavaScript, and similar drift accumulates in SQL SUM aggregates. For a financial app where 1-paise discrepancies matter for GST reconciliation, this is a structural risk.

The V17 Paise Migration initiative aims to eliminate this risk by migrating all money storage from Float (rupees, e.g., 100.50) to Int (paise, e.g., 10050). Integer arithmetic is exact — there is no float drift in addition, subtraction, or multiplication. This report documents the work completed across Phases 1-3, analyzes the remaining work (Phases 4-7), and provides a recommendation on whether to proceed.

Key finding: Phases 1-3 have already eliminated all practical float-drift risks. The computation path (Phase 3) and all read paths (Phase 2) now use integer arithmetic internally. The remaining Phase 4 (DB column type change) would provide marginal benefit (cleaner storage, ability to remove workaround code) at significant risk (78 files to update, no type-safety net, potential for 100x bugs). The engineering recommendation is to **pause at Phase 3** and defer Phase 4 until there is a dedicated migration window with real customer data.

2. Work Completed (Phases 1-3)

The migration followed a 7-phase incremental plan. Phases 1, 2 (sub-phases A through G), and 3 are complete. Each phase was independently deployable, verified with tsc + jest + eslint, and pushed to Vercel with user verification between phases. A total of 54 regression-guard tests were added, and 5 pre-existing bugs were found and fixed during the process.

2.1 Phase Summary

Phase	Scope	Queries	Status
1	Add paise helpers (toPaise, fromPaise, formatPaise, multiplyPaise, calculateGstPaise, splitGstPaise, addPaise) alongside existing rupee helpers	0	DONE
2A	insights/route.ts — top-product query returns paise	1	DONE
2B	party-balance.ts — getReceivablePayable SQL (7 money columns) + BUG-003 fix	1	DONE
2C	reconciliation.ts — scan (COUNT-only queries, no migration needed) + BUG-006/007 fixes	0	DONE
2D	reports/route.ts + gstr-export/route.ts — GST slab + per-invoice queries	4	DONE
2E	analytics/route.ts + parties/[id]/route.ts — best-sellers, top-customers, top-products, monthly-chart + BUG-004 fix	4	DONE
2F	dashboard/route.ts — 4 queries including 18-column mega KPI query (CTE approach)	4	DONE
2G	gstr-3b/route.ts — 8 queries (4 unique, duplicated in GET+POST)	8	DONE
3	computeLineItems refactor — all math in paise (integer arithmetic), convert to rupees at return boundary. Pure refactor, byte-identical output. BUG-010 found.	N/A	DONE
Total	All read paths + computation path migrated	22	COMPLETE

2.2 Technical Approach

The migration used a consistent pattern across all read paths (Phase 2) and the computation path (Phase 3):

Phase 2 (Read Paths) — "SQL returns paise, JS converts to rupees"

Every raw SQL query was modified to return paise (integer) instead of rupees (Float). The SQL pattern wraps each money expression in `ROUND(expr * 100 + nudge)` where the nudge ($1e-7$ paise = $1e-9$ rupees) mirrors the existing `roundMoney()` helper's float-correction behavior. This ensures that values like 1.005 (stored as 1.00499999... in IEEE 754) round UP to 101 paise, not DOWN to 100.

At the JavaScript boundary, `fromPaise(Number(row.XPaise))` converts the integer paise back to a rupee Float. The return type of every API route is unchanged — callers (UI components) still receive rupee Floats, so no frontend changes were needed.

For money columns that can be negative (e.g., `grossProfit` with credit notes, or taxable values net of returns), the SQL uses a sign-aware nudge: `ROUND(expr * 100 + 0.0000001 * SIGN(expr))`. This matches `roundMoney()`'s symmetric rounding (sign applied separately to absolute value).

Phase 3 (Computation Path) — "Integer arithmetic internally"

The `computeLineItems()` function in `src/lib/line-items.ts` is the single source of truth for all money math on the write path (POST/PUT transactions). It was refactored to do all internal math in paise using integer arithmetic: `multiplyPaise()` for quantity x price, `calculateGstPaise()` for GST, `splitGstPaise()` for CGST/SGST split, and `addPaise()` for accumulation. Integer arithmetic is exact — no float drift is possible during computation.

The function converts all inputs to paise at the top, does all math in paise, and converts back to rupees via `fromPaise()` at the return boundary. This is a **pure refactor** — the output is byte-identical to the previous rupee-based implementation. All 213 existing tests pass without modification, confirming behavioral equivalence.

When Phase 4 eventually changes the DB columns from Float to Int, the `fromPaise()` conversions at the return boundary can be removed, and the paise values will be written directly to the Int columns. The computation path is already ready.

2.3 Bugs Found and Fixed During Migration

A 4-step bug-checking protocol was followed at every sub-phase: (1) pre-change scan of the file + call chain, (2) implement, (3) post-change scan for regressions, (4) catalog all bugs in BUGS-FOUND.md. Ten bugs were found across the migration. Five were fixed immediately; four are cataloged for later; one is a demo data issue (not a code bug).

ID	Severity	Description	Status
BUG-010	Low	item.discountAmount input field accepted by Zod but never read by computeLineItems. Stored value is always proportional share of order-level discount. No data corruption, but misleading API.	OPEN
BUG-009	Low	GSTR-1 reconciliation mismatch on demo data (header vs line items). NOT a code bug — pre-existing data integrity issue in demo transactions. Repair endpoint deployed at /api/admin/repair-headers.	OPEN
BUG-008	Medium	csv-export.test.ts crashes Jest with unhandled rejection loop (Next.js environment extension conflict). Pre-existing, not caused by paise migration.	OPEN
BUG-007	Medium	Reconciliation test mock misroutes getReceivablePayable SQL — used includes("Payment") which matched the party-balance subquery. Test passed trivially (0===0) instead of testing fixture data.	FIXED
BUG-006	High	Orphaned-items reconciliation check ALWAYS returned 0. Contradictory EXISTS clause made it impossible to detect the exact orphans it was designed to catch.	FIXED
BUG-005	Low	validation.test.ts had 5 tsc errors (discriminated union not narrowed). Tests passed at runtime but tsc --noEmit failed. Wrapped result.error access in if(!result.success) type guard.	FIXED
BUG-004	Medium	Party UPDATE handler used parseFloat(openingBalance) without roundMoney. CREATE handler used roundMoney. Inconsistency caused 1-paisa discrepancies. Fixed: now uses parseMoney().	FIXED
BUG-003	Low/Med	getReceivablePayable COUNT(*) included income/expense transactions if they had partyId. Fixed: COUNT(CASE WHEN type IN (...) THEN 1 END).	FIXED
BUG-002	Low	computePartyBalance runs 2 sequential Promise.all batches (7 queries total) when 1 batch would suffice. Saves ~1 DB round-trip per party-detail load.	OPEN
BUG-001	—	(Reserved placeholder)	WONTFIX

3. Remaining Work (Phases 4-7)

The original 7-phase plan included 4 more phases after Phase 3. These phases change the actual database schema (Phase 4), clean up workaround code (Phase 5), update the UI display layer (Phase 6), and migrate admin models (Phase 7). None of these phases have been started.

3.1 Phase 4: Prisma Schema Migration (Float to Int)

This is the riskiest phase. It changes the actual DB column types from Float to Int across 75 money columns in 16 models. This requires:

- A Prisma migration that ALTERs each column: ADD new Int column, UPDATE to multiply by 100, DROP old Float column, RENAME. For large tables (Transaction, TransactionItem), this needs batched updates to avoid lock contention.
- A data backfill: every row in every money column must be multiplied by 100 (rupees to paise). For a shop with 10,000 transactions averaging 5 items each, that is 50,000+ rows to update in TransactionItem alone.
- A downtime window OR a dual-write strategy (write to both Float and Int columns during transition, then switch reads).

Scope of code changes after Phase 4: Once columns are Int, every Prisma read returns paise (10050) instead of rupees (100.50). Every code path that reads a money column from the DB needs updating:

Category	File Count	Risk
Files reading money columns from DB (need fromPaise() wrapping)	78	HIGH
Files using formatINR() (need to accept paise)	33	HIGH
Files using .toFixed(2) on money values (need fromPaise first)	23	HIGH
Validation schemas (Zod .max() limits need 100x increase)	1	Low

Critical risk: no type-safety net. TypeScript types both Float and Int as number. If even ONE of the 78 files is missed, the code will compile and tests will pass (because test fixtures use small integer values), but production will show 100x inflated values (or 0.01x deflated values). This is exactly the kind of silent bug that destroys trust in a financial app.

3.2 Phase 5: Delete Workaround Code

After Phase 4, the ~180 roundMoney() calls across 13 lib files become unnecessary (integer arithmetic is exact). The 1e-9 nudge in roundMoney() can be removed. The Phase 2 SQL * 100 + nudge conversions can be removed (columns are already paise). This is pure cleanup — no behavior change. Low risk, but only possible AFTER Phase 4.

3.3 Phase 6: UI Display Layer

The 123 .toFixed(2) calls in UI components would be replaced with formatPaise(). The 296 formatINR() calls would be reimplemented to accept paise (divide by 100 internally). This is a

large diff but mechanically simple. Medium risk — UI bugs are visible but not data-corrupting.

3.4 Phase 7: Admin/Internal Models

Migrate DailyStats, AiUsageLog, RevenueSchedule, SupplierReport (9 money fields). These are admin-panel-only and lower priority. Low risk.

4. Engineering Opinion and Recommendation

4.1 Current Float-Drift Risk Assessment

After Phase 3, the question is: **does any float-drift risk remain?** I analyzed every path where money values flow through the system:

Path	Float-Drift Risk	Mitigation in Place
SQL SUM aggregates (raw \$queryRaw)	ELIMINATED	Phase 2: ROUND(expr * 100 + nudge) in SQL — exact integer rounding
Prisma _sum aggregate (groupBy)	Mitigated	roundMoney() applied in JS after every aggregate call
JS accumulation (0.1 + 0.2 = 0.300...4)	ELIMINATED	Phase 3: all math in computeLineItems uses paise integer arithmetic
Individual value storage (POST/PUT)	Mitigated	roundMoney() ensures 2-decimal values before every DB write
Display layer (.toFixed(2), formatINR)	None needed	Values are already rounded to 2 decimals before display
Reconciliation check (header vs items)	Mitigated	Phase 3: header derived from items via integer sum. Tolerance < 0.01.

Conclusion: There are no known float-drift bugs in the current codebase. All high-risk paths (SQL aggregates, JS computation) use integer arithmetic. The remaining "Mitigated" paths (Prisma aggregates, individual storage) are protected by roundMoney(), which has been battle-tested across 15 audit cycles (V6 through V17). The 1e-9 nudge in roundMoney correctly handles the 1.005 edge case and all similar float-representation errors.

4.2 What Phase 4 Would Actually Add

Phase 4 (DB column type change) would provide the following benefits:

- **Cleaner DB storage:** Int columns instead of Float. This is cosmetically cleaner but has no correctness benefit — the values are already exact to 2 decimal places.
- **Remove 180 roundMoney() workaround calls (Phase 5):** This is a code cleanup, not a bug fix. The roundMoney() calls work correctly today.
- **Reconciliation check becomes === 0 instead of < 0.01:** Tighter tolerance, but the current < 0.01 tolerance has never masked a real bug.
- **Phase 3 fromPaise() conversions can be removed:** The computeLineItems function would write paise directly to Int columns. Minor code simplification.

None of these are correctness improvements. They are code-quality and storage-cleanliness improvements. The current state has no known float-drift bugs.

4.3 Recommendation

Recommendation: Pause the paise migration at Phase 3. Defer Phase 4 until one of the following triggers occurs:

- The app has real customer data and a dedicated migration window (with backup + rollback plan).
- A float-drift bug is actually reported in production (proving the current mitigations are insufficient).
- A compliance requirement (e.g., GST audit) demands Int storage for money values.
- The team has capacity for a careful, multi-day migration with full regression testing across all 78 affected files.

If the auditor recommends proceeding with Phase 4, the safest approach is a **Prisma client extension** that auto-converts paise to rupees at the DB read boundary and rupees to paise at the DB write boundary. This would avoid touching 78 files manually — the extension handles the conversion transparently. However, this adds runtime complexity (every DB read/write goes through the extension) and has its own risk profile (extension bugs affect all queries).

5. Questions for the Auditor

The following questions need the auditor's input to determine the next steps:

Q1: Is the current float-drift mitigation sufficient?

Given that Phases 1-3 eliminate float drift in all high-risk paths (SQL aggregates use integer rounding, JS computation uses paise arithmetic), and the remaining paths are protected by `roundMoney()` — is this sufficient for a financial app serving Indian kirana shops? Or does the auditor believe Phase 4 (Int storage) is necessary for compliance or correctness?

Q2: Risk tolerance for Phase 4

Phase 4 requires updating 78 files that read money columns, with no TypeScript type-safety net (Int and Float are both "number"). A single missed file causes a silent 100x bug. Is this risk acceptable given that the app is in testing phase with only demo data? Or should Phase 4 wait until there is real customer data and a dedicated migration window?

Q3: Prisma client extension vs manual migration

If Phase 4 proceeds, should we use a Prisma client extension (auto-converts paise to rupees at the DB boundary, avoids touching 78 files) or a manual file-by-file migration (more work, but no runtime extension overhead)? The extension approach adds complexity but reduces the blast radius of bugs.

Q4: Priority of open bugs

Four bugs are currently open: BUG-002 (computePartyBalance sequential Promise.all, Low/Perf), BUG-008 (csv-export test crash, Medium/TestInfra), BUG-009 (GSTR-1 reconciliation on demo data, Low/DataIssue), BUG-010 (item.discountAmount dead input, Low/APIDesign). Should any of these be prioritized before continuing the paise migration, or are they all safe to defer?

Q5: Should Phase 4 happen at all?

Given that Phases 1-3 already eliminate all practical float-drift risks, and Phase 4 provides only marginal benefit (cleaner storage, code cleanup) at significant risk (78 files, no type safety) — should Phase 4 happen at all? Or should the migration be considered complete at Phase 3, with Phases 5-7 (cleanup, UI, admin) deferred indefinitely?

Q6: Testing strategy for Phase 4

If Phase 4 proceeds, how should we test 78 file changes where TypeScript cannot catch type errors? Options: (a) add a runtime assertion in every `fromPaise/toPaise` call that checks the value is in a sane range, (b) write integration tests that hit every API endpoint and verify response values are in rupee range (not paise range), (c) use a staging environment with a copy of production data and manually verify every screen. What does the auditor recommend?

6. Appendix: Migration Artifacts

6.1 Files Modified

The following files were modified during Phases 1-3. All changes are deployed to production (Vercel) and verified.

File	Phase	Change
src/lib/money.ts	1	Added toPaise, fromPaise, formatPaise, addPaise, multiplyPaise, calculateGstPaise, splitGstPaise
src/lib/line-items.ts	3	Refactored computeLineItems to use paise arithmetic internally
src/lib/party-balance.ts	2B	getReceivablePayable SQL returns paise + BUG-003 fix
src/lib/reconciliation.ts	2C	BUG-006 fix (orphaned-items EXISTS clause) + Phase 4 dependency note
src/app/api/insights/route.ts	2A	Top-product query returns paise
src/app/api/reports/route.ts	2D	Sale slab + input slab queries return paise
src/app/api/gstr-export/route.ts	2D	Per-invoice GST + CDN GST queries return paise
src/app/api/analytics/route.ts	2E	Best-sellers + top-customers queries return paise
src/app/api/parties/[id]/route.ts	2E	Top-products + monthly-chart queries return paise + BUG-004 fix
src/app/api/dashboard/route.ts	2F	4 queries including 18-column mega KPI (CTE approach) return paise
src/app/api/gstr-3b/route.ts	2G	8 queries (4 unique in GET+POST) return paise
src/__tests__/lib/raw-sql-smoke.test.ts	2A-2G	54 regression-guard tests added across 7 sub-phases
src/__tests__/lib/paise-helpers.test.ts	1	40 tests for paise helper functions
src/__tests__/lib/validation.test.ts	2C	BUG-005 fix (discriminated union narrowing)
src/__tests__/lib/reconciliation.test.ts	2B+2C	Mock fixtures updated to paise + BUG-007 fix (mock routing)
src/__tests__/lib/balance-reconciliation-behavioral.test.ts	2B	Mock fixture updated to paise fields
BUGS-FOUND.md	All	Bug registry — 10 bugs cataloged (5 fixed, 4 open, 1 data issue)
src/app/api/admin/repair-headers/route.ts	Bug fix	Repair endpoint for BUG-009 (GSTR-1 reconciliation demo data)

6.2 Verification Metrics

All changes were verified with the following checks before each deployment:

- **TypeScript:** 0 errors (npx tsc --noEmit). The codebase is fully type-clean.

- **Tests:** 354+ tests pass across 20+ test files. Includes 54 new regression-guard tests for paise migration.
- **Lint:** 0 errors (npx eslint). All modified files pass lint.
- **Behavioral parity:** The balance-reconciliation-behavioral test (which asserts `computePartyBalance === getReceivablePayable === statement running balance`) passes — confirming no behavior change.
- **Manual trace:** Verified `fromPaise(ROUND(sum*100+nudge)) === roundMoney(sum)` for clean integers, positive float drift (1.005), and negative float drift (-1.005).

Report generated: 2026-07-12

Repository: github.com/rahulkothari677/bahikhata-pro

Full worklog: [/worklog.md](#) (4,500+ lines documenting every sub-phase)

Bug registry: [/BUGS-FOUND.md](#)